

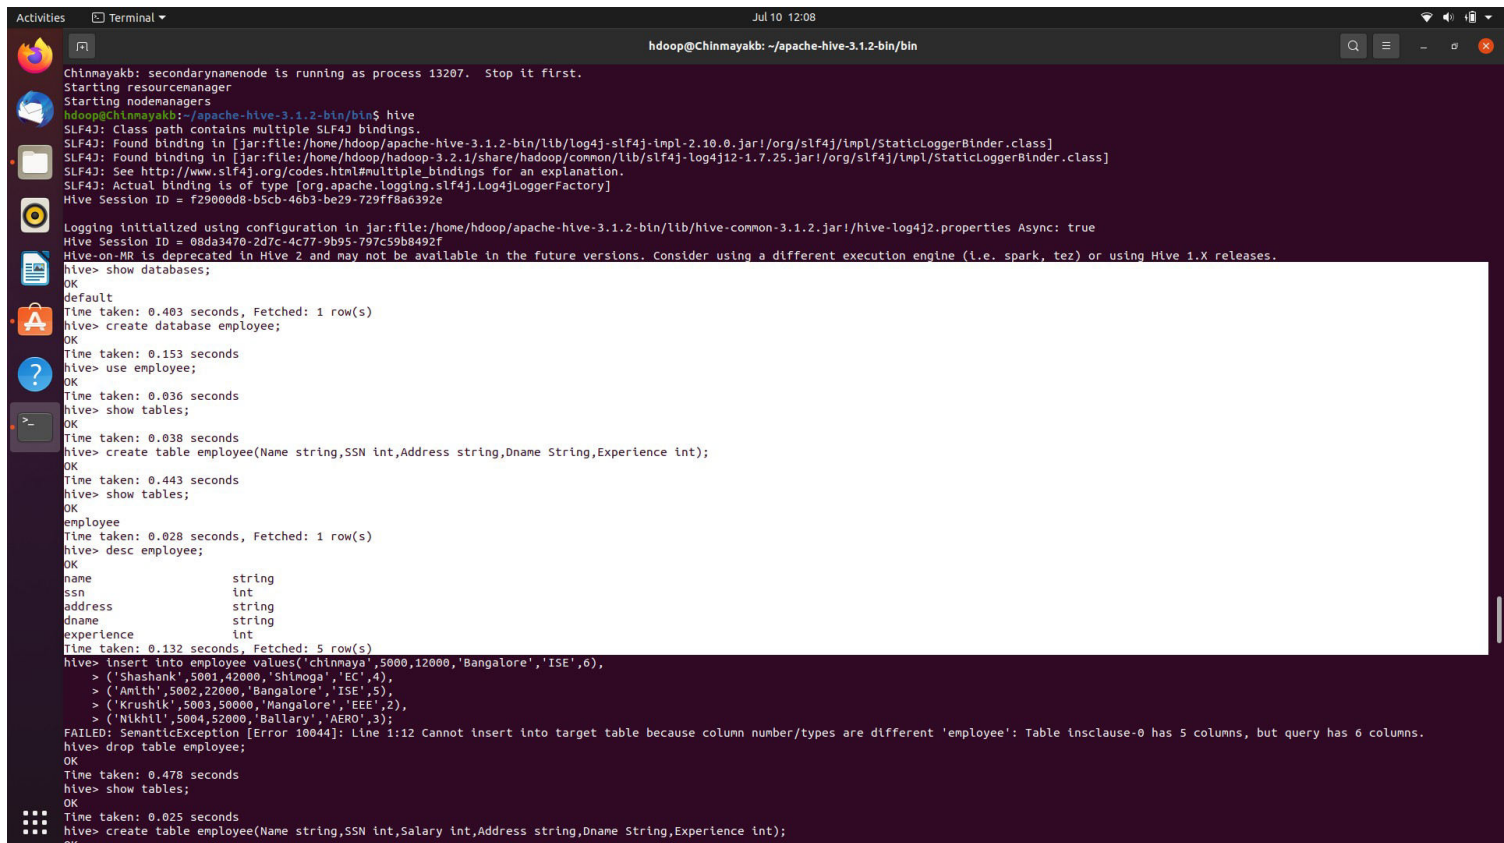
LA2

CHINMAYAKRISHNA B

1NT18IS047

Exercise2:Section 2

1) Create a database EmployeeDb and a table in it with named employee



```
hadoop@Chinnmayakb: ~/apache-hive-3.1.2-bin/bin
Chinnmayakb: secondarynamenode is running as process 13207. Stop it first.
Starting resourcemanager
Starting nodemanagers
hadoop@Chinnmayakb:~/apache-hive-3.1.2-bin/bin$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop-3.2.1/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = f2900d8-b5cb-46b3-be29-729ff8a6392e

Logging initialized using configuration in jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true
Hive Session ID = 08da3470-2d7c-4c77-9b95-797c59b8492f
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> show databases;
OK
default
Time taken: 0.403 seconds, Fetched: 1 row(s)
hive> create database employee;
OK
Time taken: 0.153 seconds
hive> use employee;
OK
Time taken: 0.036 seconds
hive> show tables;
OK
Time taken: 0.038 seconds
hive> create table employee(Name string,SSN int,Address string,Dname String,Experience int);
OK
Time taken: 0.443 seconds
hive> show tables;
OK
employee
Time taken: 0.028 seconds, Fetched: 1 row(s)
hive> desc employee;
OK
name                string
ssn                  int
address              string
dname                string
experience            int
Time taken: 0.132 seconds, Fetched: 5 row(s)
hive> insert into employee values('chinnaya',5000,12000,'Bangalore','ISE',6),
> ('Shashank',5001,42000,'Shimoga','EC',4),
> ('Amith',5002,22000,'Bangalore','ISE',5),
> ('Krushik',5003,50000,'Mangalore','EEE',2),
> ('Nikhil',5004,52000,'Ballary','AERO',3);
FAILED: SemanticException [Error 10044]: Line 1:12 Cannot insert into target table because column number/types are different 'employee': Table insclause-0 has 5 columns, but query has 6 columns.
hive> drop table employee;
OK
Time taken: 0.478 seconds
hive> show tables;
OK
Time taken: 0.025 seconds
hive> create table employee(Name string,SSN int,Salary int,Address string,Dname String,Experience int);
OK
```

2).Insert 5 Values into the Table

4).Retrieve all the employees whos salary is not less than 50000

```
Activities Terminal Jul 10 12:09 hdoop@Chinmayakb: ~/apache-hive-3.1.2-bin/bin
>
>
>
>
>
>
> show tables;
OK
employee
Time taken: 0.025 seconds, Fetched: 1 row(s)
hive> alter table employee rename to EMP;
OK
Time taken: 0.153 seconds
hive> show tables;
OK
emp
Time taken: 0.028 seconds, Fetched: 1 row(s)
hive> alter table emp change dname dept_name string;
OK
Time taken: 0.121 seconds
hive> desc emp;
OK
name                string
ssn                  int
salary               int
address              string
dept_name             string
experience             int
Time taken: 0.04 seconds, Fetched: 6 row(s)
hive> select * from emp where salary>=50000;
OK
Krushik 5003 50000 Mangalore EEE 2
Nikhil 5004 52000 Ballary AERO 3
Time taken: 0.212 seconds, Fetched: 2 row(s)
hive> select * from emp where address ='Bangalore' and experience<5;
OK
Time taken: 0.161 seconds
hive> select * from emp where address ='Bangalore' and experience>5;
OK
chinmaya 5000 12000 Bangalore ISE 6
Time taken: 0.136 seconds, Fetched: 1 row(s)
hive> create view details as (select name,dept_name from emp);
OK
Time taken: 0.159 seconds
hive> select * from details;
OK
chinmaya ISE
Shashank EC
Amith ISE
Krushik EEE
Nikhil AERO
Time taken: 0.166 seconds, Fetched: 5 row(s)
```

5).Extract all the employees who live in Bangalore but having less than 5 years of Experience

```
Activities Terminal Jul 10 12:10 hdoop@Chinmayakb: ~/apache-hive-3.1.2-bin/bin
>
>
>
>
> show tables;
OK
employee
Time taken: 0.025 seconds, Fetched: 1 row(s)
hive> alter table employee rename to EMP;
OK
Time taken: 0.153 seconds
hive> show tables;
OK
emp
Time taken: 0.028 seconds, Fetched: 1 row(s)
hive> alter table emp change dname dept_name string;
OK
Time taken: 0.121 seconds
hive> desc emp;
OK
name                string
ssn                  int
salary               int
address              string
dept_name             string
experience             int
Time taken: 0.04 seconds, Fetched: 6 row(s)
hive> select * from emp where salary>=50000;
OK
Krushik 5003 50000 Mangalore EEE 2
Nikhil 5004 52000 Ballary AERO 3
Time taken: 0.212 seconds, Fetched: 2 row(s)
hive> select * from emp where address ='Bangalore' and experience<5;
OK
Time taken: 0.161 seconds
hive> select * from emp where address ='Bangalore' and experience>5;
OK
chinmaya 5000 12000 Bangalore ISE 6
Time taken: 0.136 seconds, Fetched: 1 row(s)
hive> create view details as (select name,dept_name from emp);
OK
Time taken: 0.159 seconds
hive> select * from details;
OK
chinmaya ISE
Shashank EC
Amith ISE
Krushik EEE
Nikhil AERO
Time taken: 0.166 seconds, Fetched: 5 row(s)
hive> select name,ssn from emp group by ssn,name order by name;
Query ID = hdoop_20210710120631_8239df8f-b6af-4117-bbe4-bb1592299fbd
Total Jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
```

6).Create separate view containing Name,Dept_name of employees

```
Activities Terminal Jul 10 12:10
hdoop@Chinmayakb: ~/apache-hive-3.1.2-bin/bin

>
>
>
> show tables;
OK
employee
Time taken: 0.025 seconds, Fetched: 1 row(s)
hive> alter table employee rename to EMP;
OK
Time taken: 0.153 seconds
hive> show tables;
OK
emp
Time taken: 0.028 seconds, Fetched: 1 row(s)
hive> alter table emp change dname dept_name string;
OK
Time taken: 0.121 seconds
hive> desc emp;
OK
name                string
ssn                  int
salary              int
address              string
dept_name            string
experience            int
Time taken: 0.04 seconds, Fetched: 6 row(s)
hive> select * from emp where salary>=50000;
OK
Krushik 5003    50000    Mangalore    EEE    2
Nikhil  5004    52000    Ballary     AERO    3
Time taken: 0.212 seconds, Fetched: 2 row(s)
hive> select * from emp where address ='Bangalore' and experience<5;
OK
Time taken: 0.161 seconds
hive> select * from emp where address ='Bangalore' and experience>5;
OK
chinmaya      5000    12000    Bangalore    ISE    6
Time taken: 0.136 seconds, Fetched: 1 row(s)
hive> create view details as (select name,dept_name from emp);
OK
Time taken: 0.159 seconds
hive> select * from details;
OK
chinmaya      ISE
Shashank      EC
Amith         ISE
Krushik       EEE
Nikhil        AERO
Time taken: 0.166 seconds, Fetched: 5 row(s)
hive> select name,ssn from emp group by ssn,name order by name;
Query ID = hdoop_20210710120631_8239df8f-b6af-4117-bbe4-bb1592299fbd
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625897919713_0002, Tracking URL = http://Chinmayakb:8088/proxy/application_1625897919713_0002/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0002
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-10 12:06:38,147 Stage-1 map = 0%, reduce = 0%
2021-07-10 12:06:49,510 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.86 sec
MapReduce Total cumulative CPU time: 2 seconds 860 msec
Ended Job = job_1625897919713_0002
Launching Job 2 out of 2
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625897919713_0003, Tracking URL = http://Chinmayakb:8088/proxy/application_1625897919713_0003/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0003
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2021-07-10 12:07:00,442 Stage-2 map = 0%, reduce = 0%
2021-07-10 12:07:10,726 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 2.12 sec
2021-07-10 12:07:15,885 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 3.99 sec
MapReduce Total cumulative CPU time: 3 seconds 990 msec
Ended Job = job_1625897919713_0003
MapReduce Jobs Launched:
Stage-Stage1: Map: 1 Reduce: 1 Cumulative CPU: 2.86 sec HDFS Read: 11947 HDFS Write: 235 SUCCESS
Stage-Stage2: Map: 1 Reduce: 1 Cumulative CPU: 3.99 sec HDFS Read: 7645 HDFS Write: 211 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 850 msec
OK
Amith 5002
Krushik 5003
Nikhil 5004
Shashank 5001
chinmaya 5000
Time taken: 50.715 seconds, Fetched: 5 row(s)
hive>
```

7).Display Name and SSN and use groupby SSN and Name

```
Activities Terminal Jul 10 12:10
hdoop@Chinmayakb: ~/apache-hive-3.1.2-bin/bin

hive> select * from details;
OK
chinmaya      ISE
Shashank      EC
Amith         ISE
Krushik       EEE
Nikhil        AERO
Time taken: 0.166 seconds, Fetched: 5 row(s)
hive> select name,ssn from emp group by ssn,name order by name;
Query ID = hdoop_20210710120631_8239df8f-b6af-4117-bbe4-bb1592299fbd
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625897919713_0002, Tracking URL = http://Chinmayakb:8088/proxy/application_1625897919713_0002/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0002
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-10 12:06:38,147 Stage-1 map = 0%, reduce = 0%
2021-07-10 12:06:49,510 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.86 sec
MapReduce Total cumulative CPU time: 2 seconds 860 msec
Ended Job = job_1625897919713_0002
Launching Job 2 out of 2
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625897919713_0003, Tracking URL = http://Chinmayakb:8088/proxy/application_1625897919713_0003/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0003
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2021-07-10 12:07:00,442 Stage-2 map = 0%, reduce = 0%
2021-07-10 12:07:10,726 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 2.12 sec
2021-07-10 12:07:15,885 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 3.99 sec
MapReduce Total cumulative CPU time: 3 seconds 990 msec
Ended Job = job_1625897919713_0003
MapReduce Jobs Launched:
Stage-Stage1: Map: 1 Reduce: 1 Cumulative CPU: 2.86 sec HDFS Read: 11947 HDFS Write: 235 SUCCESS
Stage-Stage2: Map: 1 Reduce: 1 Cumulative CPU: 3.99 sec HDFS Read: 7645 HDFS Write: 211 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 850 msec
OK
Amith 5002
Krushik 5003
Nikhil 5004
Shashank 5001
chinmaya 5000
Time taken: 50.715 seconds, Fetched: 5 row(s)
hive>
```


8).Retrieve MAX salary,MIN salary and AVG salary of the employees

```
Activities Terminal Jul 10 12:11
hdoop@Chinnmayakb: ~/apache-hive-3.1.2-bin/bin

MapReduce Total cumulative CPU time: 2 seconds 860 msec
Ended Job = job_1625897919713_0002
Launching Job 2 out of 2
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625897919713_0003, Tracking URL = http://Chinnmayakb:8088/proxy/application_1625897919713_0003/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0003
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2021-07-10 12:07:00,442 Stage-2 map = 0%, reduce = 0%
2021-07-10 12:07:10,726 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.12 sec
2021-07-10 12:07:15,885 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 3.99 sec
MapReduce Total cumulative CPU time: 3 seconds 990 msec
Ended Job = job_1625897919713_0003
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.86 sec HDFS Read: 11947 HDFS Write: 235 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 3.99 sec HDFS Read: 7645 HDFS Write: 211 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 850 msec
OK
Amith 5002
Krushik 5003
Nikhil 5004
Shashank 5001
chinmaya 5000
Time taken: 50.715 seconds, Fetched: 5 row(s)
hive> select min(salary),max(salary),avg(salary) from emp;
Query ID = hdoop_20210710121111_1f39f9c4-8e66-4baa-8c5a-b71a6c33f060
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625897919713_0004, Tracking URL = http://Chinnmayakb:8088/proxy/application_1625897919713_0004/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-10 12:11:17,569 Stage-1 map = 0%, reduce = 0%
2021-07-10 12:11:21,727 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.67 sec
2021-07-10 12:11:27,908 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.43 sec
MapReduce Total cumulative CPU time: 4 seconds 430 msec
Ended Job = job_1625897919713_0004
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.43 sec HDFS Read: 17374 HDFS Write: 119 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 430 msec
OK
12000 52000 35600.0
Time taken: 17.583 seconds, Fetched: 1 row(s)
hive>
```

9).Vreate another table called department with the following fields (Dname=deptname and perform the joins.

Dno	Dname
6	ISE

```
Activities Terminal Jul 10 12:14
hdoop@Chinnmayakb: ~/apache-hive-3.1.2-bin/bin

at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:498)
at org.apache.hadoop.util.RunJar.run(RunJar.java:323)
at org.apache.hadoop.util.RunJar.main(RunJar.java:236)
FAILED: ParseException line 1:7 cannot recognize input near 'create' 'department' '(' in ddl statement
hive> create table department(dno int,dname string);
OK
Time taken: 0.12 seconds
hive> show tables;
OK
department
details
emp
Time taken: 0.03 seconds, Fetched: 3 row(s)
hive> desc department;
OK
dno int
dname string
Time taken: 0.027 seconds, Fetched: 2 row(s)
hive> insert into department values(6,'ISE'),(5,'EC'),(7,'AERO');
Query ID = hdoop_20210710121402_e316535a-0a4d-4736-adf2-5119c7f377c0
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625897919713_0005, Tracking URL = http://Chinnmayakb:8088/proxy/application_1625897919713_0005/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0005
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-10 12:14:07,051 Stage-1 map = 0%, reduce = 0%
2021-07-10 12:14:12,792 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.68 sec
2021-07-10 12:14:16,914 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.34 sec
MapReduce Total cumulative CPU time: 4 seconds 340 msec
Ended Job = job_1625897919713_0005
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/employee.db/department/.hive-staging_hive_2021-07-10_12-14-02_449_9028547012151235819-1/-ext-10000
Loading data to table employee.department
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.34 sec HDFS Read: 15628 HDFS Write: 280 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 340 msec
OK
Time taken: 15.809 seconds
hive> select * from department;
OK
6 ISE
5 EC
7 AERO
Time taken: 0.111 seconds, Fetched: 3 row(s)
hive>
```

10).Left outer Join

```
Activities Terminal Jul 10 12:16
hadoop@Chinnmayakb: ~/apache-hive-3.1.2-bin/bin

Starting Job = job_1625897919713_0005, Tracking URL = http://chinnmayakb:8088/proxy/application_1625897919713_0005/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0005
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-10 12:14:07,651 Stage-1 map = 0%, reduce = 0%
2021-07-10 12:14:12,792 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.68 sec
2021-07-10 12:14:16,914 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.34 sec
MapReduce Total cumulative CPU time: 4 seconds 340 msec
Ended Job = job_1625897919713_0005
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/employee.db/department/.hive-staging_hive_2021-07-10_12-14-02_449_9028547812151235819-1/-ext-10000
Loading data to table employee.department
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.34 sec HDFS Read: 15628 HDFS Write: 280 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 340 msec
OK
Time taken: 15.809 seconds
hive> select * from department;
OK
6 ISE
5 EC
7 AERO
Time taken: 0.111 seconds, Fetched: 3 row(s)
hive> select * from emp left outer join department where dname=dept_name;
Query ID = hadoop_20210710121555_13bc8842-0de7-4185-becf-f28133639a3d
Total jobs = 1

SLF4J: Found binding in [jar:file:/home/hadoop/hadoop-3.2.1/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2021-07-10 12:16:01 Dump the side-table for tag: 1 with group count: 3 into file: file:/tmp/hive/java/hadoop/f29000d8-b5cb-46b3-be29-729ffa6392e/hive_2021-07-10_12-15-55_658_5060518414468933463-1/-lo
cal-10004/HashTable-Stage-3/MapJoin-mapfile01---.hashtable2021-07-10 12:16:01 Uploaded 1 file to: file:/tmp/hive/java/hadoop/f29000d8-b5cb-46b3-be29-729ffa6392e/hive_2021-07-10_12-15-55_658_50605184144
68933463-1/-local-10004/HashTable-Stage-3/MapJoin-mapfile01---.hashtable (329 bytes)
Execution completed successfully
MapReduce task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625897919713_0006, Tracking URL = http://chinnmayakb:8088/proxy/application_1625897919713_0006/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0006
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2021-07-10 12:16:08,263 Stage-3 map = 0%, reduce = 0%
2021-07-10 12:16:12,381 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 2.45 sec
MapReduce Total cumulative CPU time: 2 seconds 450 msec
Ended Job = job_1625897919713_0006
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 2.45 sec HDFS Read: 10419 HDFS Write: 294 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 450 msec
OK
chinnaya 5000 12000 Bangalore ISE 6 6 ISE
Shashank 5001 42000 Shinoga EC 4 5 EC
Amith 5002 22000 Bangalore ISE 5 6 ISE
Nikhil 5004 52000 Ballary AERO 3 7 AERO
Time taken: 24.036 seconds, Fetched: 4 row(s)
hive>
```

11).Right Outer Join

```
Activities Terminal Jul 10 12:17
hadoop@Chinnmayakb: ~/apache-hive-3.1.2-bin/bin

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2021-07-10 12:16:01 Dump the side-table for tag: 1 with group count: 3 into file: file:/tmp/hive/java/hadoop/f29000d8-b5cb-46b3-be29-729ffa6392e/hive_2021-07-10_12-15-55_658_5060518414468933463-1/-lo
cal-10004/HashTable-Stage-3/MapJoin-mapfile01---.hashtable2021-07-10 12:16:01 Uploaded 1 file to: file:/tmp/hive/java/hadoop/f29000d8-b5cb-46b3-be29-729ffa6392e/hive_2021-07-10_12-15-55_658_50605184144
68933463-1/-local-10004/HashTable-Stage-3/MapJoin-mapfile01---.hashtable (329 bytes)
Execution completed successfully
MapReduce task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625897919713_0006, Tracking URL = http://chinnmayakb:8088/proxy/application_1625897919713_0006/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0006
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2021-07-10 12:16:08,263 Stage-3 map = 0%, reduce = 0%
2021-07-10 12:16:12,381 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 2.45 sec
MapReduce Total cumulative CPU time: 2 seconds 450 msec
Ended Job = job_1625897919713_0006
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 2.45 sec HDFS Read: 10419 HDFS Write: 294 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 450 msec
OK
chinnaya 5000 12000 Bangalore ISE 6 6 ISE
Shashank 5001 42000 Shinoga EC 4 5 EC
Amith 5002 22000 Bangalore ISE 5 6 ISE
Nikhil 5004 52000 Ballary AERO 3 7 AERO
Time taken: 24.036 seconds, Fetched: 4 row(s)
hive> select * from emp right outer join department where dname=dept_name;
Query ID = hadoop_20210710121645_a828ba60-0338-48e8-9029-88773660df32
Total jobs = 1

SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop-3.2.1/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2021-07-10 12:16:50 Dump the side-table for tag: 1 with group count: 3 into file: file:/tmp/hive/java/hadoop/f29000d8-b5cb-46b3-be29-729ffa6392e/hive_2021-07-10_12-16-45_561_8658834088026560490-1/-lo
cal-10004/HashTable-Stage-3/MapJoin-mapfile11---.hashtable
2021-07-10 12:16:50 End of local task; Time Taken: 1.151 sec.
Execution completed successfully
MapReduce task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625897919713_0007, Tracking URL = http://chinnmayakb:8088/proxy/application_1625897919713_0007/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0007
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2021-07-10 12:16:57,442 Stage-3 map = 0%, reduce = 0%
2021-07-10 12:17:02,587 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 2.5 sec
MapReduce Total cumulative CPU time: 2 seconds 500 msec
Ended Job = job_1625897919713_0007
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 2.5 sec HDFS Read: 10432 HDFS Write: 294 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 500 msec
OK
chinnaya 5000 12000 Bangalore ISE 6 6 ISE
Shashank 5001 42000 Shinoga EC 4 5 EC
Amith 5002 22000 Bangalore ISE 5 6 ISE
Nikhil 5004 52000 Ballary AERO 3 7 AERO
Time taken: 18.452 seconds, Fetched: 4 row(s)
hive>
```

12).Full Outer Join

```
Activities Terminal Jul 10 12:18
hadoop@Chinnmayakb: ~/apache-hive-3.1.2-bin/bin

Total jobs = 1
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop-3.2.1/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2021-07-10 12:16:50 Dump the side-table for tag: 1 with group count: 3 into file: file:/tmp/hive/java/hadoop/f29000d8-b5cb-46b3-be29-729ff8a6392e/hive_2021-07-10_12-16-45_561_8658834088026560490-1/-local-10004/HashTable-Stage-3/MapJoin-mapfile11--.hashtable
2021-07-10 12:16:50 End of local task; Time Taken: 1.151 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625897919713_0007, Tracking URL = http://Chinnmayakb:8088/proxy/application_1625897919713_0007/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0007
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2021-07-10 12:16:57,442 Stage-3 map = 0%, reduce = 0%
2021-07-10 12:17:02,587 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 2.5 sec
MapReduce Total cumulative CPU time: 2 seconds 500 msec
Ended Job = job_1625897919713_0007
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 2.5 sec HDFS Read: 10432 HDFS Write: 294 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 500 msec
OK
chinnaya 5000 12000 Bangalore ISE 6 6 ISE
Shashank 5001 42000 Shimoga EC 4 5 EC
Amith 5002 22000 Bangalore ISE 5 6 ISE
Nikhil 5004 52000 Ballary AERO 3 7 AERO
Time taken: 18.452 seconds, Fetched: 4 row(s)
hive> select * from emp full outer join department where dname=dept_name;
Query ID = hadoop_20210710121720_f5c9d77c-9493-43c6-8f3f-4fa714303d1e
Total jobs = 1
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
2021-07-10 12:17:25 Dump the side-table for tag: 1 with group count: 3 into file: file:/tmp/hive/java/hadoop/f29000d8-b5cb-46b3-be29-729ff8a6392e/hive_2021-07-10_12-17-20_112_5810221896728672370-1/-local-10004/HashTable-Stage-3/MapJoin-mapfile21--.hashtable2021-07-10 12:17:25 Uploaded 1 file to: file:/tmp/hive/java/hadoop/f29000d8-b5cb-46b3-be29-729ff8a6392e/hive_2021-07-10_12-17-20_112_5810221896728672370-1/-local-10004/HashTable-Stage-3/MapJoin-mapfile21--.hashtable (329 bytes)
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625897919713_0008, Tracking URL = http://Chinnmayakb:8088/proxy/application_1625897919713_0008/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0008
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2021-07-10 12:17:31,105 Stage-3 map = 0%, reduce = 0%
2021-07-10 12:17:36,203 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 2.61 sec
MapReduce Total cumulative CPU time: 2 seconds 610 msec
Ended Job = job_1625897919713_0008
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 2.61 sec HDFS Read: 10432 HDFS Write: 294 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 610 msec
OK
chinnaya 5000 12000 Bangalore ISE 6 6 ISE
Shashank 5001 42000 Shimoga EC 4 5 EC
Amith 5002 22000 Bangalore ISE 5 6 ISE
Nikhil 5004 52000 Ballary AERO 3 7 AERO
Time taken: 17.194 seconds, Fetched: 4 row(s)
hive>
```

CODE:

```
hive> show databases;
```

OK

default

Time taken: 0.403 seconds, Fetched: 1 row(s)

```
hive> create database employee;
```

OK

Time taken: 0.153 seconds

```
hive> use employee;
```

OK

Time taken: 0.036 seconds

```
hive> show tables;
```

OK

Time taken: 0.038 seconds

```
hive> create table employee(Name string,SSN int,Address string,Dname  
String,Experience int);
```

OK

Time taken: 0.443 seconds

```
hive> show tables;
```

OK

employee

Time taken: 0.028 seconds, Fetched: 1 row(s)

```
hive> desc employee;
```

OK

name	string
------	--------

ssn	int
-----	-----

address	string
---------	--------

dname	string
-------	--------

experience	int
------------	-----

Time taken: 0.132 seconds, Fetched: 5 row(s)

```
hive> insert into employee values('chinmaya',5000,12000,'Bangalore','ISE',6),
```

```
> ('Shashank',5001,42000,'Shimoga','EC',4),
```

```
> ('Amith',5002,22000,'Bangalore','ISE',5),
```

```
> ('Krushik',5003,50000,'Mangalore','EEE',2),
```

```
> ('Nikhil',5004,52000,'Ballary','AERO',3);
```

FAILED: SemanticException [Error 10044]: Line 1:12 Cannot insert into target table because column number/types are different 'employee': Table insclause-0 has 5 columns, but query has 6 columns.

```
hive> drop table employee;
```

OK

Time taken: 0.478 seconds

```
hive> show tables;
```

OK

Time taken: 0.025 seconds


```
hive> create table employee(Name string,SSN int,Salary int,Address string,Dname String,Experience int);
```

OK

Time taken: 0.109 seconds

```
hive> insert into employee values('chinmaya',5000,12000,'Bangalore','ISE',6),  
> ('Shashank',5001,42000,'Shimoga','EC',4),  
> ('Amith',5002,22000,'Bangalore','ISE',5),  
> ('Krushik',5003,50000,'Mangalore','EEE',2),  
> ('Nikhil',5004,52000,'Ballary','AERO',3);
```

Query ID = hdoop_20210710120100_c46890b9-092e-45ec-84e1-4315c11e11c4

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1625897919713_0001, Tracking URL = http://Chinmayakb:8088/proxy/application_1625897919713_0001/

Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill

job_1625897919713_0001

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2021-07-10 12:01:09,670 Stage-1 map = 0%, reduce = 0%

2021-07-10 12:01:15,889 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.19 sec

2021-07-10 12:01:21,007 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.75 sec

MapReduce Total cumulative CPU time: 3 seconds 750 msec

Ended Job = job_1625897919713_0001

Stage-4 is selected by condition resolver.

Stage-3 is filtered out by condition resolver.

Stage-5 is filtered out by condition resolver.

Moving data to directory

hdfs://127.0.0.1:9000/user/hive/warehouse/employee.db/employee/.hive-staging_hive_2021-07-10_12-01-00_533_3418192329285162114-1/-ext-10000

Loading data to table employee.employee

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.75 sec HDFS Read: 20983

HDFS Write: 671 SUCCESS

Total MapReduce CPU Time Spent: 3 seconds 750 msec

OK

Time taken: 23.141 seconds

```
hive> select * from employee;
```

```
OK
chinmaya  5000 12000Bangalore  ISE  6
Shashank  5001 42000Shimoga    EC   4
Amith     5002 22000Bangalore  ISE  5
Krushik   5003 50000Mangalore  EEE  2
Nikhil    5004 52000Ballary    AERO  3
Time taken: 0.206 seconds, Fetched: 5 row(s)
```

```
> show tables;
```

```
OK
```

```
employee
```

```
Time taken: 0.025 seconds, Fetched: 1 row(s)
```

```
hive> alter table employee rename to EMP;
```

```
OK
```

```
Time taken: 0.153 seconds
```

```
hive> show tables;
```

```
OK
```

```
emp
```

```
Time taken: 0.028 seconds, Fetched: 1 row(s)
```

```
hive> alter table emp change dname dept_name string;
```

```
OK
```

```
Time taken: 0.121 seconds
```

```
hive> desc emp;
```

```
OK
```

```
name          string
ssn            int
salary        int
address       string
dept_name     string
experience     int
```

```
Time taken: 0.04 seconds, Fetched: 6 row(s)
```

```
hive> select * from emp where salary>=50000;
```

```
OK
```

```
Krushik   5003 50000Mangalore  EEE  2
Nikhil    5004 52000Ballary    AERO  3
```

```
Time taken: 0.212 seconds, Fetched: 2 row(s)
```

```
hive> select * from emp where address = 'Bangalore' and experience<5;
```

```
OK
```

```
Time taken: 0.161 seconds
```

```
hive> select * from emp where address = 'Bangalore' and experience>5;
```

```
OK
```

```
chinmaya  5000 12000Bangalore  ISE  6
```

```
Time taken: 0.136 seconds, Fetched: 1 row(s)
```

```
hive> create view details as (select name,dept_name from emp);
```

```
OK
```

Time taken: 0.159 seconds

hive> select * from details;

OK

chinmaya ISE

Shashank EC

Amith ISE

Krushik EEE

Nikhil AERO

Time taken: 0.166 seconds, Fetched: 5 row(s)

hive> select name,ssn from emp group by ssn,name order by name;

Query ID = hdoop_20210710120631_8239df8f-b6af-4117-bbe4-bb1592299fbd

Total jobs = 2

Launching Job 1 out of 2

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1625897919713_0002, Tracking URL = http://Chinmayakb:8088/proxy/application_1625897919713_0002/

Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill
job_1625897919713_0002

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2021-07-10 12:06:38,147 Stage-1 map = 0%, reduce = 0%

2021-07-10 12:06:43,301 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.4 sec

2021-07-10 12:06:49,510 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.86 sec

MapReduce Total cumulative CPU time: 2 seconds 860 msec

Ended Job = job_1625897919713_0002

Launching Job 2 out of 2

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1625897919713_0003, Tracking URL = http://Chinmayakb:8088/proxy/application_1625897919713_0003/

Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill
job_1625897919713_0003

Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1

2021-07-10 12:07:00,442 Stage-2 map = 0%, reduce = 0%
2021-07-10 12:07:10,726 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.12 sec
2021-07-10 12:07:15,885 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 3.99 sec
MapReduce Total cumulative CPU time: 3 seconds 990 msec
Ended Job = job_1625897919713_0003
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.86 sec HDFS Read: 11947 HDFS Write: 235 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 3.99 sec HDFS Read: 7645 HDFS Write: 211 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 850 msec
OK
Amith 5002
Krushik 5003
Nikhil 5004
Shashank 5001
chinmaya 5000
Time taken: 50.715 seconds, Fetched: 5 row(s)
hive> select min(salary),max(salary),avg(salary) from emp;
Query ID = hdoop_20210710121111_1f39f9c4-8e66-4baa-8c5a-b71a6c33f060
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
set mapreduce.job.reduces=<number>
Starting Job = job_1625897919713_0004, Tracking URL = http://Chinmayakb:8088/proxy/application_1625897919713_0004/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-10 12:11:17,569 Stage-1 map = 0%, reduce = 0%
2021-07-10 12:11:21,727 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.67 sec
2021-07-10 12:11:27,908 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.43 sec
MapReduce Total cumulative CPU time: 4 seconds 430 msec
Ended Job = job_1625897919713_0004
MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.43 sec HDFS Read: 17374
HDFS Write: 119 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 430 msec
OK
120005200035600.0
Time taken: 17.583 seconds, Fetched: 1 row(s)

hive> create table department(dno int,dname string);
OK

Time taken: 0.12 seconds

hive> show tables;

OK

department

details

emp

Time taken: 0.03 seconds, Fetched: 3 row(s)

hive> desc department;

OK

dno	int
-----	-----

dname	string
-------	--------

Time taken: 0.027 seconds, Fetched: 2 row(s)

hive> insert into department values(6,'ISE'),(5,'EC'),(7,'AERO');

Query ID = hdoop_20210710121402_e316535a-0a4d-4736-adf2-5119c7f377c0

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1625897919713_0005, Tracking URL = http://Chinmayakb:8088/proxy/application_1625897919713_0005/

Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill

job_1625897919713_0005

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2021-07-10 12:14:07,651 Stage-1 map = 0%, reduce = 0%

2021-07-10 12:14:12,792 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.68 sec

2021-07-10 12:14:16,914 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.34 sec

MapReduce Total cumulative CPU time: 4 seconds 340 msec

Ended Job = job_1625897919713_0005

Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory
hdfs://127.0.0.1:9000/user/hive/warehouse/employee.db/department/.hive-staging_hive_2021-07-10_12-14-02_449_9028547012151235819-1/-ext-10000
Loading data to table employee.department
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.34 sec HDFS Read: 15628
HDFS Write: 280 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 340 msec
OK
Time taken: 15.809 seconds
hive> select * from department;
OK
6 ISE
5 EC
7 AERO
Time taken: 0.111 seconds, Fetched: 3 row(s)
hive> select * from emp left outer join department where dname=dept_name;
Query ID = hdoop_20210710121555_13bc8842-0de7-4185-becf-f28133639a3d
Total jobs = 1

SLF4J: Found binding in
[jar:file:/home/hadoop/hadoop-3.2.1/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2021-07-10 12:16:01 Dump the side-table for tag: 1 with group count: 3 into file:
file:/tmp/hive/java/hadoop/f29000d8-b5cb-46b3-be29-729ff8a6392e/hive_2021-07-10_12-15-55_658_5060518414468933463-1/-local-10004/HashTable-Stage-3/MapJoin-mapfile01--.hashtable
2021-07-10 12:16:01 Uploaded 1 File to:
file:/tmp/hive/java/hadoop/f29000d8-b5cb-46b3-be29-729ff8a6392e/hive_2021-07-10_12-15-55_658_5060518414468933463-1/-local-10004/HashTable-Stage-3/MapJoin-mapfile01--.hashtable (329 bytes)
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625897919713_0006, Tracking URL = http://Chinmayakb:8088/proxy/application_1625897919713_0006/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0006
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2021-07-10 12:16:08,263 Stage-3 map = 0%, reduce = 0%

2021-07-10 12:16:12,381 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 2.45 sec

MapReduce Total cumulative CPU time: 2 seconds 450 msec

Ended Job = job_1625897919713_0006

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 2.45 sec HDFS Read: 10419 HDFS

Write: 294 SUCCESS

Total MapReduce CPU Time Spent: 2 seconds 450 msec

OK

chinmaya	5000	12000	Bangalore	ISE	6	6	ISE
Shashank	5001	42000	Shimoga	EC	4	5	EC
Amith	5002	22000	Bangalore	ISE	5	6	ISE
Nikhil	5004	52000	Ballary	AERO	3	7	AERO

Time taken: 24.036 seconds, Fetched: 4 row(s)

hive> select * from emp right outer join department where dname=dept_name;

Query ID = hdoop_20210710121645_a828ba60-0338-48e8-9029-88773660df32

Total jobs = 1

SLF4J: Found binding in [jar:file:/home/hdoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in

[jar:file:/home/hdoop/hadoop-3.2.1/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

2021-07-10 12:16:50 Dump the side-table for tag: 1 with group count: 3 into file: file:/tmp/hive/java/hdoop/f29000d8-b5cb-46b3-be29-729ff8a6392e/hive_2021-07-10_12-16-45_561_8658834088026560490-1/-local-10004/HashTable-Stage-3/

MapJoin-mapfile11--.hashtable

2021-07-10 12:16:50 End of local task; Time Taken: 1.151 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1625897919713_0007, Tracking URL = http://Chinmayakb:8088/proxy/application_1625897919713_0007/

Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0007

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2021-07-10 12:16:57,442 Stage-3 map = 0%, reduce = 0%

2021-07-10 12:17:02,587 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 2.5 sec

MapReduce Total cumulative CPU time: 2 seconds 500 msec

Ended Job = job_1625897919713_0007

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 2.5 sec HDFS Read: 10432 HDFS Write: 294 SUCCESS

Total MapReduce CPU Time Spent: 2 seconds 500 msec

OK

chinmaya	5000	12000	Bangalore	ISE	6	6	ISE
Shashank	5001	42000	Shimoga	EC	4	5	EC
Amith	5002	22000	Bangalore	ISE	5	6	ISE
Nikhil	5004	52000	Ballary	AERO	3	7	AERO

Time taken: 18.452 seconds, Fetched: 4 row(s)

hive> select * from emp full outer join department where dname=dept_name;

Query ID = hdoop_20210710121720_f5c9d77c-9493-43c6-8f3f-4fa714303d1e

Total jobs = 1

SLF4J: Found binding in [jar:file:/home/hdoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]

2021-07-10 12:17:25 Dump the side-table for tag: 1 with group count: 3 into file:

file:/tmp/hive/java/hdoop/f29000d8-b5cb-46b3-be29-729ff8a6392e/hive_2021-07-10_12-17-20_112_5810221896728672370-1/-local-10004/HashTable-Stage-3/

MapJoin-mapfile21--.hashtable2021-07-10 12:17:25 Uploaded 1 File to:

file:/tmp/hive/java/hdoop/f29000d8-b5cb-46b3-be29-729ff8a6392e/hive_2021-07-10_12-17-20_112_5810221896728672370-1/-local-10004/HashTable-Stage-3/

MapJoin-mapfile21--.hashtable (329 bytes)

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1625897919713_0008, Tracking URL = http://Chinmayakb:8088/proxy/application_1625897919713_0008/

Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625897919713_0008

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2021-07-10 12:17:31,105 Stage-3 map = 0%, reduce = 0%

2021-07-10 12:17:36,203 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 2.61 sec

MapReduce Total cumulative CPU time: 2 seconds 610 msec

Ended Job = job_1625897919713_0008

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 2.61 sec HDFS Read: 10432 HDFS

Write: 294 SUCCESS

Total MapReduce CPU Time Spent: 2 seconds 610 msec

OK

chinmaya	5000	12000	Bangalore	ISE	6	6	ISE
Shashank	5001	42000	Shimoga	EC	4	5	EC
Amith	5002	22000	Bangalore	ISE	5	6	ISE
Nikhil	5004	52000	Ballary	AERO	3	7	AERO

Time taken: 17.194 seconds, Fetched: 4 row(s)

hive>

